



华南理工大学

South China University of Technology

The Experiment Report of Machine Learning

SCHOOL: SCHOOL OF SOFTWARE ENGINEERING

SUBJECT: SOFTWARE ENGINEERING

Author:
Siyuan Xiao

Supervisor:
Qingyao Wu

Student ID:
201721045886

Grade:
Undergraduate

December 14, 2017

Logistic Regression, Linear Classification and Stochastic Gradient Descent

Abstract—Logistic Regression is a method using regression techniques to achieve classification. Stochastic Gradient Descent is an advanced way to solve optimizing problems, and is very popular in the field of machine learning. This experiment combines these two.

I. INTRODUCTION

This report will talk about the whole experiment I have made on Logistic Regression and Linear classification, which are based on Stochastic Gradient Descent. Its content is organized as follow:

- 1) Section II contains the experiment steps.
- 2) Section III contains the code for the two experiments.
- 3) Section IV makes conclusion for the experiment result.

II. METHODS AND THEORY

Experiment uses a9a of LIBSVM Data, including 32561/16281(testing) samples and each sample has 123/123 (testing) features. Please download the training set and validation set.

Then the experiment will be performed by the following steps:

- 1) Download the dataset to local host machine.
- 2) Load dataset into memory.
- 3) Split dataset into training set and validation set.
- 4) Create and fill necessary data structures according to different optimizing methods.
- 5) Write functions for calculating loss and gradient (different in regression and classification).
- 6) Set parameters for different optimizing methods (learning rate and the number of iterations).
- 7) Initiate weights (using normal distribution).
- 8) Calculate the gradient and update weights according to the optimizing methods.

- 9) Switch to another optimizing method and run again.
- 10) Draw plot for experiment result.

III. EXPERIMENT

Here I placed the code for the two experiments:

1) Logistic Regression:

```

1. # created by Swain, 2017-12-14, 13:25
2.
3. import math
4. import numpy
5. import matplotlib.pyplot as plot
6. from numpy import random
7. from sklearn.datasets import
   load_svmlight_file
8.
9. #load a9a dataset
10. #training dataset
11. data = load_svmlight_file('./a9a')
12. X_train = data[0].toarray()
13. y_train = data[1]
14. data = load_svmlight_file('./a9a.t')
15. X_vali = data[0].toarray()
16. y_vali = data[1]
17.
18. #complete the martrix
19. X_vali = numpy.column_stack((X_vali,
   numpy.zeros((X_vali.shape[0])))
20.
21. #add a constant-bias-column to X
22. X_train = numpy.column_stack((X_train,
   numpy.ones((X_train.shape[0])))
23. X_vali = numpy.column_stack((X_vali,
   numpy.ones((X_vali.shape[0])))
24.
25. #create weight array with initial values in
   normal distribution
26. d = X_train.shape[1]
27. W_init = numpy.random.normal(size=d)
28.
29. #define loss function
30. def loss(X, W, y, _lambda):
31.     y_predict = numpy.dot(X, W)
32.     return numpy.sum(numpy.log(1 +
   numpy.exp(-y * y_predict))) / X.shape[0] +
   _lambda / 2 * numpy.dot(W, W.T)
33.
34. #define gradient function

```

```

35. def grad(X, W, y, _lambda):
36.     y_predict = numpy.dot(X, W)
37.     return numpy.dot(((1 - y) / (1 + numpy.exp(y
    * y_predict))), X) / X.shape[0] + W * _lambda
38.
39. #shuffle the array
40. def shuffle_array(X_train):
41.     randomlist =
    numpy.arange(X_train.shape[0])
42.     numpy.random.shuffle(randomlist)
43.     X_random = X_train[randomlist]
44.     y_random = y_train[randomlist]
45.     return X_random, y_random
46.
47. #get the training instance and label in current
    batch
48. def
    get_Batch(runs, X_random, y_random, batch_si
    ze, shape):
49.     if k == runs - 1:
50.         X_batch = X_random[k * batch_size :
    shape + 1]
51.         y_batch = y_random[k * batch_size :
    shape + 1]
52.     else:
53.         X_batch = X_random[k * batch_size :
    (k+1) * batch_size]
54.         y_batch = y_random[k * batch_size :
    (k+1) * batch_size]
55.     return X_batch, y_batch
56.
57. #parameters: learning rate and #iteration
58. lr = 0.05
59. epoch = 5
60. batch_size = 128
61. runs = math.ceil(X_train.shape[0] /
    float(batch_size))
62. iteration = epoch * runs
63.
64. _lambda = 0.01
65.
66. #NAG/AdaDelta
67. gamma = 0.8
68.
69. #RMSprop/AdaDelta/Adam
70. epsilon = numpy.e**(-8)
71.
72. #Adam
73. beta1 = 0.9
74. beta2 = 0.999
75.
76. #used to save results
77. loss_train = []
78. loss_vali = []
79.
80. nmethods = 4
81. #use different optimizing method for each i
82. for i in range(0, nmethods):
83.     #reset W

```

```

84.     W = W_init
85.     loss_train.append(numpy.zeros(iteration))
86.     loss_vali.append(numpy.zeros(iteration))
87.
88.     #NAG
89.     vt = numpy.zeros(X_train.shape[1])
90.
91.     #RMSprop/AdaDelta
92.     g2 = 0
93.
94.     #AdaDelta
95.     w2 = 0
96.     RMS_g = 0
97.     RMS_w = 0
98.     w_delta = numpy.zeros(X_train.shape[1])
99.
100.    #Adam
101.    mt = numpy.zeros(X_train.shape[1])
102.    nt = 0
103.
104.    for j in range(0, epoch):
105.        X_random, y_random =
        shuffle_array(X_train)
106.        for k in range(0, runs):
107.            #get a batch of training data
108.            X_batch, y_batch =
            get_Batch(runs, X_random, y_random, batch_si
            ze, X_train.shape[0])
109.
110.            #calculate gradient
111.            #NAG
112.            if i == 0:
113.                G = grad(X_batch, W - vt *
                gamma, y_batch, _lambda)
114.            #others
115.            else:
116.                G = grad(X_batch, W, y_batch,
                _lambda)
117.
118.            #calculate loss on both training and
            validation datasets
119.            loss_train[i][j * runs + k] =
            loss(X_batch, W, y_batch, _lambda)
120.            loss_vali[i][j * runs + k] =
            loss(X_vali, W, y_vali, _lambda)
121.
122.            #update weight according to
            optimizing methods
123.            #NAG
124.            if (i == 0):
125.                vt = vt * gamma + G * lr
126.                W = W - vt
127.            #RMSprop
128.            elif (i == 1):
129.                g2 = g2 * 0.9 + numpy.dot(G, G.T)
                * 0.1
130.                W = W - G * (lr / math.sqrt(g2 +
                epsilon))
131.            elif (i == 2):

```

```

132.         g2 = g2 * gamma +
numpy.dot(G,G.T) * (1-gamma)
133.         RMS_g = math.sqrt(g2 + epsilon)
134.         W = W - G *(RMS_w / RMS_g)
135.         w_delta = G *(- lr / RMS_g)
136.         w2 = w2 * gamma +
numpy.dot(w_delta, w_delta.T) * (1-gamma)
137.         RMS_w = math.sqrt(w2 + epsilon)
138.     else:
139.         mt = mt * beta1 + G * (1-beta1)
140.         nt = nt * beta2 +
numpy.dot(G,G.T) * (1-beta2)
141.         hat_m = mt * (1/(1-beta1))
142.         hat_n = nt * (1/(1-beta2))
143.         W = W - hat_m *
(lr/(math.sqrt(hat_n)+epsilon))
144.
145.     names = ['NAG', 'RMSprop', 'AdaDelta',
'Adam']
146.     i = 0
147.
148.     plot.plot(loss_train[i], label="training loss")
149.     plot.plot(loss_val[i], label="validation loss")
150.     plot.legend()
151.     plot.xlabel("Iteration")
152.     plot.ylabel("Validation Loss")
153.     plot.title('Logistic Regression + ' +
names[i])
154.     plot.show()

```

2) Linear Classification:

The only difference from the former is the loss function and its corresponding gradient function.

```

1.  #define loss function (Hinge loss)
2.  def loss(X, W, y, _lambda):
3.      y_predict = numpy.dot(X,W)
4.      diff = numpy.ones(y.shape[0]) - y *
y_predict
5.      diff[diff < 0] = 0
6.      W_0 = W.copy()
7.      W_0[len(W) - 1] = 0
8.      return numpy.sum(diff) / X.shape[0] +
numpy.dot(W_0,W_0.T) / 2 * _lambda


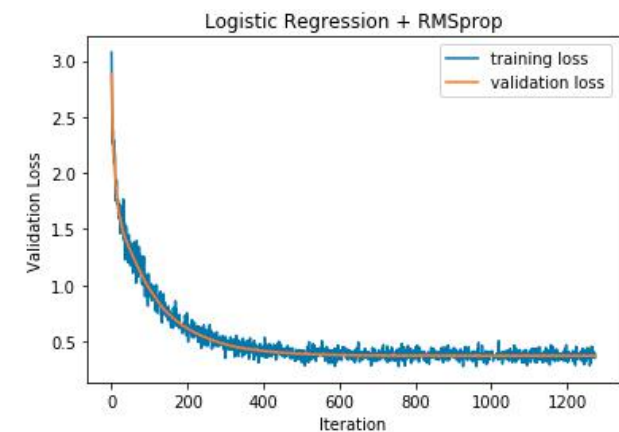
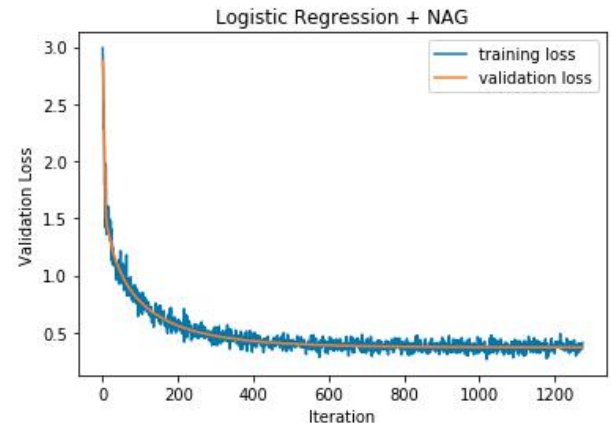
9.  #define gradient function
10. def grad(X, W, y, _lambda):
11.     y_predict = numpy.dot(X,W)
12.     diff = numpy.ones(y.shape[0]) - y *
y_predict
13.     y_ = y.copy()
14.     y_[diff <= 0] = 0
15.     W_0 = W.copy()
16.     W_0[len(W) - 1] = 0
17.     return -numpy.dot(y_,X) / X.shape[0] + W_0
* _lambda



```

IV. CONCLUSION

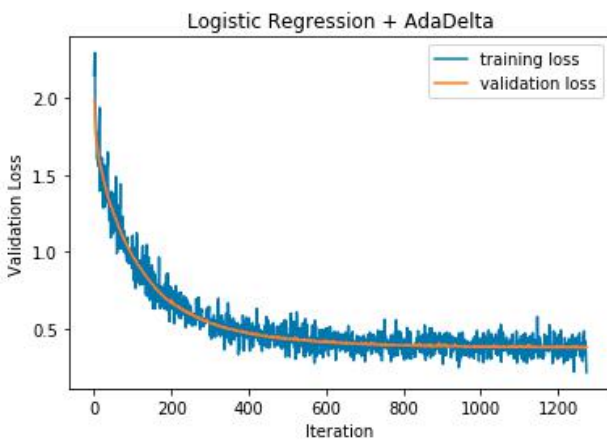


Here is the experiment result gained:

1) Logistic Regression





2) Linear Classification



Then we can draw a conclusion according to the two experiments:

- 1) Using batches to train the model is faster and more effective.
- 2) The 4 optimizing methods show slight ference on loss in these two experiments.